

Security, Auth, and Data Protection

Identity, RBAC, ownership, isolation, middleware, and sensitive data.

Eurskem AI · Agentic Workflow Platform

Generated 2026-08-24

Code revision 8bd16ca4a2a3

Derived from the live repository

Source of truth. Inventories and symbol tables are generated from the current checkout. Narrative explains observed structures and does not replace executable contracts.

Security model

Authentication establishes a user; route dependencies enforce roles. Resource handlers verify ownership. Middleware adds request context, size limits, trusted hosts, CORS, and Redis rate limits. Retrieval scope is enforced in query compilation, not prompts. Protected files are proxied through authenticated routes.

- Secrets come from configuration and are not returned in manifests.
- Entity tokenization can replace sensitive values before model calls.
- External actions are constrained and audited.
- Production rate limiting fails closed when Redis is unavailable.
- Security filters are separate from self-query metadata filters.

File	Responsibility
<code>app/api/auth.py</code>	Auth module.
<code>app/retrieval/__init__.py</code>	Hybrid RAG retrieval pipeline.
<code>app/retrieval/compressor.py</code>	Stage 4 of the retrieval pipeline: contextual compression.
<code>app/retrieval/context.py</code>	Parent/sentence-window/contextual expansion and final context assembly.
<code>app/retrieval/contracts.py</code>	Small interfaces for retrieval, fusion, reranking and context expansion.
<code>app/retrieval/filters.py</code>	Typed metadata filters and the non-removable authorization-filter clauses.
<code>app/retrieval/fusion.py</code>	Reciprocal Rank Fusion for independently executed retrieval result sets.
<code>app/retrieval/hybrid_search.py</code>	Stage 2 of the retrieval pipeline: hybrid search with metadata pre-filter.
<code>app/retrieval/models.py</code>	Pydantic models for the retrieval pipeline.
<code>app/retrieval/presets.py</code>	Editable Retrieval Playground starting configurations.
<code>app/retrieval/query_transform.py</code>	Rewrite, multi-query, decomposition, HyDE and validated self-query.
<code>app/retrieval/query_understanding.py</code>	Stage 1 of the retrieval pipeline: query understanding.
<code>app/retrieval/reranker.py</code>	Stage 3 of the retrieval pipeline: LLM-as-judge reranking.
<code>app/retrieval/retriever.py</code>	The retrieval pipeline orchestrator.
<code>app/retrieval/service.py</code>	Canonical retrieval path for workflows, RAG, Playground, MCP and Eval.
<code>app/retrieval/strategies.py</code>	Replaceable dense, sparse and hybrid Weaviate retrieval strategies.
<code>app/retrieval/weaviate_client.py</code>	Weaviate client and schema definition.
<code>app/security/__init__.py</code>	Python module.
<code>app/security/audit.py</code>	Audit log writer + reader.
<code>app/security/dependencies.py</code>	Dependencies module.
<code>app/security/entity_ner.py</code>	Lazy-loaded spaCy NER safety net — PERSON/ORG spans not already known.
<code>app/security/entity_protection_errors.py</code>	Fail-closed exception hierarchy for confidential entity protection.

File	Responsibility
<code>app/security/entity_registry.py</code>	Explicit per-scope entity registry — the primary detection mechanism.
<code>app/security/entity_tokenizer.py</code>	Entity tokenization orchestration — the Phase 1 core guarantee.
<code>app/security/entity_vault.py</code>	Encrypted, reversible placeholder<->original-value vault.
<code>app/security/guardrails.py</code>	Deterministic input/output guardrails for workflow execution.
<code>app/security/jwt_handler.py</code>	Jwt handler module.
<code>app/security/llm_policy.py</code>	Data classification + OPA policy enforcement for direct provider calls.
<code>app/security/middleware.py</code>	HTTP request identity, limits, security headers, and Redis rate limiting.
<code>app/security/passwords.py</code>	Argon2 password hashing helpers for local production authentication.
<code>app/security/presidio_client.py</code>	Presidio Analyzer client — the PII/entity detection safety-net tier.
<code>app/security/rbac.py</code>	Rbac module.
<code>app/security/sso_stub.py</code>	Microsoft SSO stub.
<code>app/security/users.py</code>	Mongo-backed local users used by the IONOS production deployment.
<code>app/storage/__init__.py</code>	Python module.
<code>app/storage/minio_client.py</code>	S3-compatible object storage client.
<code>ui/src/components/auth/LoginPage.test.tsx</code>	React/TypeScript module.
<code>ui/src/components/auth/LoginPage.tsx</code>	Informational banner (e.g.

Method	Path	Auth	Handler	Purpose
GET	<code>/api/</code>	public	<code>app/api/workflows.py:146</code> <code>list_workflow_root</code>	List the workflows.
GET	<code>/api/audit/session/{session_id}</code>	consultant	<code>app/api/audit.py:15</code> <code>get_session_audit</code>	Return the session audit.
POST	<code>/api/builder/assist/rules</code>	consultant	<code>app/api/builder.py:1032</code> <code>assist_rules</code>	Turn a described business rule into deterministic configuration (§36).
POST	<code>/api/builder/assist/schema</code>	consultant	<code>app/api/builder.py:913</code> <code>assist_schema</code>	Propose a structured-output schema from a description (§37).
GET	<code>/api/builder/email/connect/{provider}</code>	consultant	<code>app/api/email_oauth.py:146</code> <code>connect</code>	Starts the authorization-code flow: stash a one-time state (+ PKCE verifier for Gmail), redirect the user's browser to the provider.
GET	<code>/api/builder/email/connections</code>	consultant	<code>app/api/builder.py:1300</code> <code>email_connections</code>	Configured mailboxes for the Email node's connection picker.
DELETE	<code>/api/builder/email/connections/{connection_id}</code>	consultant	<code>app/api/email_oauth.py:303</code> <code>disconnect</code>	Compute the disconnect.
PATCH	<code>/api/builder/email/connections/{connection_id}</code>	consultant	<code>app/api/email_oauth.py:274</code> <code>update_connection</code>	The only lever to actually let a connection send — every connection (OAuth-established or static) starts refused; this is the explicit second step, same reasoning as every other write-capable connection in this platform (MCP, F&O).

Method	Path	Auth	Handler	Purpose
GET	/api/builder/email/oauth/callback/{provider}	public	app/api/email_oauth.py:178 oauth_callback	The provider redirects the user's browser back here after consent.
GET	/api/builder/integrations/connect/{provider}	consultant	app/api/integration_oauth.py:137 connect	Starts the authorization-code flow: stash a one-time state (+ PKCE verifier for Google), redirect the user's browser to the provider.
GET	/api/builder/integrations/connections	consultant	app/api/integration_oauth.py:249 list_connections	List the connections.
DELETE	/api/builder/integrations/connections/{connection_id}	consultant	app/api/integration_oauth.py:270 disconnect	Compute the disconnect.
GET	/api/builder/integrations/connections/{connection_id}/download/{file_id}	consultant	app/api/integration_oauth.py:359 download_file	Streams one file's real bytes straight to the browser — a plain navigation (an ` <a href="">` click, not a fetch), authenticated by the same HttpOnly cookie the OAuth popup flow relies on.
GET	/api/builder/integrations/connections/{connection_id}/files	consultant	app/api/integration_oauth.py:295 browse_files	Live folder/search browsing for the node config panel's file picker — not workflow execution, just a thin proxy through the service for whichever operation the presence of `query` selects.
GET	/api/builder/integrations/connections/{connection_id}/path/{file_id}	consultant	app/api/integration_oauth.py:327 browse_path	Breadcrumb chain for the file browser: walk `parent_id` up from `file_id` via repeated lookups, bounded so a cyclic/misbehaving provider response can never hang this request.
GET	/api/builder/integrations/oauth/callback/{provider}	public	app/api/integration_oauth.py:169 oauth_callback	The provider redirects the user's browser back here after consent.
GET	/api/builder/mcp/servers	consultant	app/api/builder.py:1124 mcp_servers	Configured MCP servers for the Builder's connection panel.
GET	/api/builder/mcp/servers/{server_id}/health	consultant	app/api/builder.py:1176 mcp_health	Live connection check, for the panel's Test Connection button.
GET	/api/builder/mcp/servers/{server_id}/tools	consultant	app/api/builder.py:1142 mcp_tools	Discover what a server can do (§5).
POST	/api/builder/mcp/test-tool	consultant	app/api/builder.py:1207 mcp_test_tool	Run one MCP tool with literal arguments (§21).
POST	/api/builder/node-test	consultant	app/api/builder.py:358 node_test	Run one node against sample data (§21).
GET	/api/builder/operators	consultant	app/api/builder.py:112 operators	Which operators apply to which field type, and how to render each one.
POST	/api/builder/output-contract	consultant	app/api/builder.py:203 output_contract	What each step guarantees, as typed dotted paths (§40).
POST	/api/builder/schema-preview	consultant	app/api/builder.py:296 schema_preview	Compile visual schema rows and show what the model will be held to.

Method	Path	Auth	Handler	Purpose
POST	/api/builder/simulate	consultant	app/api/builder.py:507 simulate	Run a workflow in memory and return a per-step trace (\$23, \$24, \$47).
GET	/api/candidates/{run_id}	consultant	app/api/candidates.py:268 list_candidates	Every acquired evidence PDF for a run, plus discovered candidates.
GET	/api/candidates/{run_id}/documents/{document_id}/download	consultant	app/api/candidates.py:70 download_acquired_research_document	Authorize against owned run provenance, then issue a short-lived PDF URL.
POST	/api/candidates/{run_id}/verify-claim	consultant	app/api/candidates.py:341 verify_claim_endpoint	Run the gpt-5.6-sol secondary check for one internal-evidence record.
POST	/api/chat-conversations/resolve	consultant	app/api/chat_conversations.py:71 resolve_conversation	No summary declared.
GET	/api/chat-conversations/{conversation_id}	consultant	app/api/chat_conversations.py:90 get_conversation	No summary declared.
POST	/api/chat-conversations/{conversation_id}/messages	consultant	app/api/chat_conversations.py:108 append_message	No summary declared.
PUT	/api/chat-conversations/{conversation_id}/messages/{message_id}	consultant	app/api/chat_conversations.py:124 replace_message	No summary declared.
GET	/api/chat-workflows	consultant	app/api/chat_workflows.py:547 list_private_chat_workflows	No summary declared.
GET	/api/chat-workflows/adapters/by-name/{workflow_name}	consultant	app/api/chat_workflows.py:602 get_builder_chat_execution_adapter	Build a universal, execution-only Chat wrapper around a saved workflow.
POST	/api/chat-workflows/from-existing	consultant	app/api/chat_workflows.py:667 copy_existing_chat_workflow	No summary declared.
POST	/api/chat-workflows/generate	consultant	app/api/chat_workflows.py:688 generate_private_chat_workflow	No summary declared.
POST	/api/chat-workflows/import	consultant	app/api/chat_workflows.py:652 import_private_chat_workflow	No summary declared.
POST	/api/chat-workflows/presets/deep-research	consultant	app/api/chat_workflows.py:562 ensure_deep_research_chat_workflow	Return the owner's vetted Deep Research workflow for one Knowledge binding.
POST	/api/chat-workflows/presets/general	consultant	app/api/chat_workflows.py:585 ensure_general_chat_workflow	Return one stable owner-scoped workflow for general Chat requests.
DELETE	/api/chat-workflows/{chat_workflow_id}	consultant		No summary declared.

Method	Path	Auth	Handler	Purpose
			app/api/ chat_workflows.py:726 archive_private_chat_workflow	
GET	/api/chat-workflows/ {chat_workflow_id}	consultant	app/api/ chat_workflows.py:637 get_private_chat_workflow	No summary declared.
POST	/api/chat-workflows/ {chat_workflow_id}/ request-publication	consultant	app/api/ chat_workflows.py:738 request_private_workflow_public ation	No summary declared.
GET	/api/chat-workspace/ experiences	consultant	app/api/ chat_workspace.py:135 list_workspace_experiences	No summary declared.
POST	/api/chat-workspace/plan	consultant	app/api/ chat_workspace.py:211 plan_chat_workspace	No summary declared.
POST	/api/chat-workspace/ prepare	consultant	app/api/ chat_workspace.py:220 prepare_chat_workspace	No summary declared.
POST	/api/chat-workspace/ workflow-labels	consultant	app/api/ chat_workspace.py:143 chat_workflow_labels	Return Chat-only business labels without renaming saved workflows.
GET	/api/cost-admin/budgets	admin	app/api/ cost_admin.py:417 get_budgets	Return the budgets.
PUT	/api/cost-admin/budgets/ global	admin	app/api/ cost_admin.py:467 set_global_budget	Set the global budget.
PUT	/api/cost-admin/budgets/ session/{session_id}	admin	app/api/ cost_admin.py:494 set_session_budget	Set the session budget.
GET	/api/cost-admin/cache- summary	admin	app/api/ cost_admin.py:342 cache_summary	Compute the cache summary.
GET	/api/cost-admin/infra- allocations	admin	app/api/ cost_admin.py:225 list_infra_allocations	List the infra allocations.
PUT	/api/cost-admin/infra- allocations/{model}	admin	app/api/ cost_admin.py:256 set_infra_allocation	Set the infra allocation.
GET	/api/cost-admin/overview	admin	app/api/ cost_admin.py:531 cost_overview	Compute the cost overview.
GET	/api/cost-admin/pricing	admin	app/api/cost_admin.py:77 list_pricing	Direct/local catalog pricing (editable overrides layered on MODEL_PRICING defaults) plus a live, read-only slice of OpenRouter's own per-model pricing — OpenRouter's ~400-500 models are priced by OpenRouter itself and reported per-call via usage.cost (LedgerEntry.cost_source == "provider_reported"),

Method	Path	Auth	Handler	Purpose
DELETE	/api/cost-admin/pricing/{model}	admin	app/api/cost_admin.py:165 clear_pricing_override	Clear the pricing override.
PUT	/api/cost-admin/pricing/{model}	admin	app/api/cost_admin.py:125 set_pricing_override	Set the pricing override.
GET	/api/cost/run/{run_id}	authenticated	app/api/cost.py:14 run_cost	Run the cost.
GET	/api/cost/session/{session_id}	authenticated	app/api/cost.py:34 session_cost	Compute the session cost.
DELETE	/api/entity-registry	consultant	app/api/entity_registry.py:144 delete_entity	Delete the entity.
GET	/api/entity-registry	consultant	app/api/entity_registry.py:86 list_entities	List the entities.
POST	/api/entity-registry	consultant	app/api/entity_registry.py:112 register_entity	Register the entity.
GET	/api/entity-registry/entity-types	public	app/api/entity_registry.py:76 list_entity_types	List the entity types.
GET	/api/eval/golden-set	authenticated	app/api/eval.py:116 list_golden_set	List the golden set.
GET	/api/eval/history	authenticated	app/api/eval.py:211 history	Compute the history.
POST	/api/eval/run	authenticated	app/api/eval.py:138 run	Run the result.
POST	/api/eval/score-output	authenticated	app/api/eval.py:71 score_output	Score the output.
POST	/api/eval/workflow-compare	authenticated	app/api/eval.py:255 workflow_compare	Run the same golden set of real customer messages through the flagship triage workflow once per candidate model, and compare which one actually reaches the right business outcome — not a judged quality score, since there's a deterministic right answer (intent/complexity/route/review) to check against
GET	/api/eval/workflow-golden-set	authenticated	app/api/eval.py:234 list_workflow_golden_set	List the workflow golden set.
GET	/api/files	public	app/api/workflows.py:1345 get_file	Stream a file from object storage by key (used by the Output Viewer).
GET	/api/inspect/chunks	authenticated	app/api/inspect.py:192 list_chunks	List ingested chunks grouped by source document.
POST	/api/inspect/retrieve	authenticated	app/api/inspect.py:345 retrieve_preview	Run a query through hybrid search + rerank and show the ranked chunks.
GET	/api/knowledge/collections	authenticated	app/api/knowledge.py:110 list_collections	List the collections.

Method	Path	Auth	Handler	Purpose
POST	/api/knowledge/collections	authenticated	app/api/knowledge.py:92 create_collection	Create the collection.
GET	/api/knowledge/collections/{collection_id}	authenticated	app/api/knowledge.py:124 get_collection	Return the collection.
GET	/api/knowledge/collections/{collection_id}/documents	authenticated	app/api/knowledge.py:669 list_documents	List the documents.
GET	/api/knowledge/collections/{collection_id}/indexes	authenticated	app/api/knowledge.py:346 list_indexes	List the indexes.
POST	/api/knowledge/collections/{collection_id}/indexes	authenticated	app/api/knowledge.py:323 create_index	Create the index.
POST	/api/knowledge/collections/{collection_id}/indexes/{index_id}/activate	authenticated	app/api/knowledge.py:363 activate_index	Compute the activate index.
POST	/api/knowledge/collections/{collection_id}/ingestions	authenticated	app/api/knowledge.py:386 start_ingestion	Start the ingestion.
GET	/api/knowledge/doc-types	authenticated	app/api/knowledge.py:276 doc_types	Document types a Collection may declare.
GET	/api/knowledge/documents/{document_id}	authenticated	app/api/knowledge.py:686 get_document	Return the document.
GET	/api/knowledge/documents/{document_id}/source-url	authenticated	app/api/knowledge.py:705 document_source_url	Compute the document source url.
GET	/api/knowledge/embedding-models	authenticated	app/api/knowledge.py:288 embedding_models	Embedding models selectable when building an Index.
GET	/api/knowledge/ingestion-presets	authenticated	app/api/knowledge.py:264 ingestion_presets	Compute the ingestion presets.
GET	/api/knowledge/ingestions	authenticated	app/api/knowledge.py:606 list_ingestions	List the ingestions.
GET	/api/knowledge/ingestions/{ingestion_job_id}	authenticated	app/api/knowledge.py:622 get_ingestion	Return the ingestion.
POST	/api/knowledge/ingestions/{ingestion_job_id}/cancel	authenticated	app/api/knowledge.py:641 cancel_ingestion	Cancel the ingestion.
GET	/api/knowledge/profiles	authenticated	app/api/knowledge.py:185 list_profiles	List the profiles.
POST	/api/knowledge/profiles	authenticated	app/api/knowledge.py:164 create_profile	Create the profile.
POST	/api/knowledge/profiles/defaults	authenticated	app/api/knowledge.py:222 create_default_profiles	Create the default profiles.

Method	Path	Auth	Handler	Purpose
GET	/api/knowledge/profiles/{profile_id}	authenticated	app/api/knowledge.py:201 get_profile	Return the profile.
GET	/api/llm/models	consultant	app/api/llm_providers.py:44 list_models	Return public capability metadata without URLs or credentials.
GET	/api/llm/models/openrouter	consultant	app/api/llm_providers.py:90 list_openrouter_models	Live OpenRouter model catalog (via OpenRouter's own `GET /v1/models`), cached with a short TTL — see app/llm/openrouter_catalog.py.
POST	/api/llm/models/{model}/probe	admin	app/api/llm_providers.py:112 probe_model	Probe an enabled local endpoint; never return its private address.
GET	/api/node-types	admin	app/api/workflows.py:152 list_node_types	The Workflow Builder UI calls this to populate the palette and forms.
POST	/api/node-types/ask	consultant	app/api/node_types_chat.py:153 ask_about_node_types	Compute the ask about node types.
POST	/api/node-types/draft-code	consultant	app/api/node_types_chat.py:592 draft_code	Draft the code.
POST	/api/node-types/draft-instructions	consultant	app/api/node_types_chat.py:507 draft_instructions	Draft the instructions.
POST	/api/node-types/draft-prompt	consultant	app/api/node_types_chat.py:383 draft_prompt	Draft the prompt.
GET	/api/node-types/{node_type}/models	public	app/api/workflows.py:158 allowed_models	Returns the default allowed model list for a node type.
GET	/api/prompt-templates	consultant	app/api/prompt_templates.py:39 list_prompt_templates	No summary declared.
POST	/api/prompt-templates	consultant	app/api/prompt_templates.py:45 create_prompt_template	No summary declared.
DELETE	/api/prompt-templates/{template_id}	consultant	app/api/prompt_templates.py:85 delete_prompt_template	No summary declared.
PUT	/api/prompt-templates/{template_id}	consultant	app/api/prompt_templates.py:50 update_prompt_template	No summary declared.
POST	/api/prompt-templates/{template_id}/duplicate	consultant	app/api/prompt_templates.py:60 duplicate_prompt_template	No summary declared.
PUT	/api/prompt-templates/{template_id}/favorite	consultant	app/api/prompt_templates.py:75 favorite_prompt_template	No summary declared.
POST	/api/proposals/coverage	consultant	app/api/proposals.py:114 coverage	Compute the coverage.
GET		consultant		Compute the review run.

Method	Path	Auth	Handler	Purpose
	/api/proposals/runs/{run_id}/review		app/api/proposals.py:129 review_run	
GET	/api/proposals/{proposal_id}/approvals	consultant	app/api/proposals.py:466 list_approvals	List the approvals.
POST	/api/proposals/{proposal_id}/approvals	consultant	app/api/proposals.py:395 request_approval	Compute the request approval.
POST	/api/proposals/{proposal_id}/approvals/{approval_id}/decision	consultant	app/api/proposals.py:435 decide_approval	Compute the decide approval.
POST	/api/proposals/{proposal_id}/concept-alternatives	consultant	app/api/proposals.py:341 concept_alternatives	Compute the concept alternatives.
POST	/api/proposals/{proposal_id}/evidence/verify	consultant	app/api/proposals.py:569 run_proposal_evidence_factory	Verify already-fetched full text and return a citable evidence package.
POST	/api/proposals/{proposal_id}/horizon-evaluation	consultant	app/api/proposals.py:507 horizon_evaluation	Compute the horizon evaluation.
POST	/api/proposals/{proposal_id}/render	consultant	app/api/proposals.py:659 render_proposal	Render a proposal and persist both sanitised HTML and final PDF.
POST	/api/proposals/{proposal_id}/render/docx	consultant	app/api/proposals.py:696 render_proposal_docx	Render a proposal as editable DOCX and persist its sanitised source.
GET	/api/proposals/{proposal_id}/sources/versions	consultant	app/api/proposals.py:225 list_source_versions	List the source versions.
POST	/api/proposals/{proposal_id}/sources/{source_id}/versions	consultant	app/api/proposals.py:188 register_source_version	Register the source version.
POST	/api/proposals/{proposal_id}/verify-claims	consultant	app/api/proposals.py:261 verify_claims	Verify the claims.
GET	/api/rag-agents	authenticated	app/api/rag_agents.py:69 list_agents	List the agents.
POST	/api/rag-agents	authenticated	app/api/rag_agents.py:48 create_agent	Create the agent.
GET	/api/rag-agents/{rag_agent_id}	authenticated	app/api/rag_agents.py:87 get_agent	Return the agent.
POST	/api/rag-agents/{rag_agent_id}/query	authenticated	app/api/rag_agents.py:121 query_agent	Query the agent.
GET	/api/research/skills	consultant	app/api/research.py:10 list_scientific_skills	List the scientific skills.
POST	/api/retrieval/compare	authenticated	app/api/retrieval.py:191 compare	Compare the result.
GET	/api/retrieval/presets	authenticated	app/api/retrieval.py:20 presets	Return editable starting configurations; saving creates a profile.

Method	Path	Auth	Handler	Purpose
POST	/api/retrieval/search	authenticated	app/api/retrieval.py:158 search	Search the result.
GET	/api/retrieval/traces	authenticated	app/api/retrieval.py:239 list_traces	List the traces.
GET	/api/retrieval/traces/{retrieval_request_id}	authenticated	app/api/retrieval.py:255 get_trace	Return the trace.
GET	/api/retrieval/traces/{retrieval_request_id}/chunks/{chunk_id}/context	authenticated	app/api/retrieval.py:276 get_trace_chunk_context	Return cited and adjacent same-document passages from one scoped trace.
GET	/api/runs/mine	consultant	app/api/runs.py:50 my_runs	Return a bounded, newest-first list of the caller's global workflow runs.
DELETE	/api/runs/mine/{run_id}	consultant	app/api/runs.py:758 delete_run_endpoint	Delete the run endpoint.
GET	/api/runs/mine/{run_id}	consultant	app/api/runs.py:67 my_run_detail	Compute the my run detail.
GET	/api/runs/mine/{run_id}/chat	consultant	app/api/run_chat.py:227 get_run_chat	Return the run chat.
POST	/api/runs/mine/{run_id}/chat	consultant	app/api/run_chat.py:248 ask_about_run	Compute the ask about run.
POST	/api/runs/mine/{run_id}/pause	consultant	app/api/runs.py:454 pause_run	Ask a running workflow to pause at its next node boundary.
GET	/api/runs/mine/{run_id}/pending-gate	consultant	app/api/runs.py:101 pending_gate	The human-review gate a paused run is currently waiting on, if any.
POST	/api/runs/mine/{run_id}/resume	consultant	app/api/runs.py:500 resume_paused_run	Resume a run paused by the run-history "pause" action.
POST	/api/runs/mine/{source_run_id}/restart	consultant	app/api/runs.py:580 restart_run	Create a brand-new run of the same workflow and original inputs.
POST	/api/runs/mine/{source_run_id}/retry	consultant	app/api/runs.py:253 retry_failed_run	Create a new attempt that reuses every completed node checkpoint.
GET	/api/runs/{run_id}/events	consultant	app/api/workflows.py:1275 stream_run_events	Authenticated Server-Sent Events stream for one workflow run.
POST	/api/workflow-input-files	consultant	app/api/workflow_files.py:72 upload_workflow_input_files	Store files and return stable references suitable for workflow state.
GET	/api/workflow-input-files/capabilities	consultant	app/api/workflow_files.py:53 capabilities	Return the authoritative UI picker limits and accepted extensions.
GET	/api/workflow-input-files/content	consultant	app/api/workflow_files.py:171 download_workflow_input_file	Download one uploaded input after enforcing the authenticated scope.
POST	/api/workflow-input-files/extract	consultant	app/api/workflow_files.py:148 extract_workflow_input_file	Convert one uploaded document into plain text for the HITL editor.
GET	/api/workflows	public	app/api/workflows.py:901 list_workflows	Library view in the UI calls this to render the workflow collection.

Method	Path	Auth	Handler	Purpose
POST	/api/workflows/autofix	consultant	app/api/workflows.py:420 autofix_workflow	Best-effort repair of a workflow's preflight errors: deterministic patches first (free, predictable), then an LLM repair loop for whatever remains.
POST	/api/workflows/autofix-node	consultant	app/api/workflows.py:501 autofix_node	Repair one selected node while rejecting every unrelated YAML change.
GET	/api/workflows/by-name/{name}	public	app/api/workflows.py:984 get_workflow	Builder calls this to load an existing workflow for editing.
GET	/api/workflows/chat-catalog	public	app/api/workflows.py:951 list_chat_workflows	Return the small published-workflow catalog used by Business Chat.
POST	/api/workflows/generate	consultant	app/api/workflow_generation.py:940 generate_workflow_endpoint	Generate the workflow endpoint.
GET	/api/workflows/preflight-stats	admin	app/api/workflows.py:620 get_preflight_stats	Platform-wide preflight reliability rollup across every /validate, /generate, and /autofix call — success rate, and which error codes recur most before any fix is attempted.
POST	/api/workflows/run	consultant	app/api/workflows.py:640 run	Run the result.
POST	/api/workflows/save	public	app/api/workflows.py:1061 save_workflow	Builder calls this on Save.
POST	/api/workflows/subprocess-callback/{token}	public	app/api/workflows.py:858 subprocess_callback	Deliver a Subprocess child's terminal result to its waiting parent.
POST	/api/workflows/validate	consultant	app/api/workflows.py:231 validate_workflow	Validate a workflow without running a node or consuming LLM tokens.
DELETE	/api/workflows/{name}	consultant	app/api/workflows.py:1082 delete_workflow	Permanently removes a workflow's YAML together with its autosave draft and full version history — there is no undo.
GET	/api/workflows/{name}/detail	public	app/api/workflows.py:997 get_workflow_detail	Workflow Library details panel calls this to refresh 'library' and 'readiness' for one workflow on demand (e.g.
DELETE	/api/workflows/{name}/draft	public	app/api/workflows.py:1180 delete_workflow_draft	Delete the workflow draft.
GET	/api/workflows/{name}/draft	public	app/api/workflows.py:1170 get_workflow_draft	Builder calls this on load to offer 'Resume draft' for a newer autosave.
PUT	/api/workflows/{name}/draft	public	app/api/workflows.py:1159 save_workflow_draft	Autosave a recoverable draft.
GET	/api/workflows/{name}/stats	consultant	app/api/workflows.py:1032 get_workflow_stats	Library "Runs and performance" tab.
GET	/api/workflows/{name}/versions	public	app/api/workflows.py:1192 list_workflow_versions	List immutable versions without rewriting or preloading their YAML.

Method	Path	Auth	Handler	Purpose
GET	/api/workflows/{name}/versions/{version_id}	public	app/api/workflows.py:1204 get_workflow_version	Return the workflow version.
POST	/api/workflows/{name}/versions/{version_id}/restore	public	app/api/workflows.py:1224 restore_workflow_version	Restoring an old version is Save-equivalent: it must pass the same zero-token preflight gate before it can become the live workflow again.
POST	/api/workflows/{run_id}/resume	consultant	app/api/workflows.py:764 resume	Resume doesn't need services for execution — the cached compiled graph in <code>_PAUSED_GRAPHS</code> already has services bound into every node instance.
POST	/auth/logout	public	app/api/auth.py:148 logout	Clear the auth cookie.
GET	/auth/me	consultant	app/api/auth.py:134 me	Rehydrate identity from the auth cookie after a page refresh.
POST	/auth/token	public	app/api/auth.py:76 login	Compute the login.
GET	/health	public	app/api/health.py:264 health	Return live dependency status without failing the liveness request.
GET	/ready	public	app/api/health.py:272 ready	Return 503 until every required runtime dependency is reachable.